

Unit 1 Introduction

Unit Outcomes. Here you will learn

- to define a distributed system (DS) and describe and justify the characteristics of all DSs
- to classify DSs by various important aspects, appreciate the challenges in accurately specifying the behaviour of a DS
- about some common ways to distribute responsibility among networked nodes
- about network abstractions available to DS developers
- about basic tools for deploying and managing DSs

Further Reading: CDK2005 1, 2.3, 3, 4

Contents

1 Characteristics of a DS

- Definition of Distributed System
- Some common DSs
- Concurrency
- No global clock
- Independent failures

2 Classification of DSs

- Purpose
- Scalability
- Security
- Heterogeneity support
- Openness
- Transparency

3 Network abstraction

- Stacking abstraction layers
- Role of protocols
- Internet protocols for DS developers
- Web layer
- Middleware
- Synchronisation of communicating nodes
- Multicast

4 Remote administration

- Why remote administration?
- Remote shell
- Remote execution of GUI applications

Definition of Distributed System

A Distributed System (DS) is

- a software system
 - deployed on networked multiple computers (nodes)
 - making these nodes cooperate towards a shared goal
 - nodes communicate only by exchanging messages

Some common DSs


- phone networks (phones and exchanges)
- email (email clients, SMTP servers and mail-box servers)
- World-Wide Web (web servers and browsers)
- Internet search engines (high-performance clusters)
- enterprise systems (database servers and clients)
- folding@home, seti@home
- bank systems
- Torrent P2P
- IRC
-

Characteristics of a DS

Concurrency

- messages arrive concurrently, are processed concurrently

example:




- multi-threading inside a node usually necessary
- access to internal data must be synchronised

Characteristics of a DS

Concurrency

- messages arrive concurrently, are processed concurrently


example:



- multi-threading inside a node usually necessary
- access to internal data must be synchronised

No global clock


- impossible to 100% synchronise clocks among nodes
- approximate synchronisation using Cristian protocol:
 - estimating message delays:



- using radio or satellite signals:
- more generally: a node cannot fully know others' state

Independent failures

- failures occur in components of DS
- how to communicate failures?
 - when node dies completely
 - only by lack of response
- when network or its part fails silently
 - slow or broken?
 - nodes may be stranded, DS fragmented




Classification of DSs

Purpose

- information transfer over distance
- resource sharing • coordination
- performance enhancement
 - processing of large data
 - complex scientific computation
- protection against node failure by duplicating resources

Scalability

- how does a DS cope with increasing load?
- load = eg count of nodes, requests, users, ...
- perfectly scalable DS:
 - can be adapted to take any load
 - without modifying protocols or sw
 - cost of hw grows proportionally: $O(n)$
 - decrease in performance logarithmic: $O(\log(n))$ time increase for basic ops



Security — attacks

- *security* = level of resistance to attacks
- *attack* = attempt to make a DS behave in unintended ways
- typical goals of attacks:
 - unauthorised access to information
 - unauthorised alteration of data
 - total system break-down

Security — attack techniques


- typical techniques used in attacks:
 - deploy a *rogue node* to the DS
 - *spy* on messages in the DS
 - *modify* messages in the DS
 - *flood* the DS with legitimate but meaningless messages
 - send *malicious mobile code*

Security — levels and solutions

- various levels of security in DS by combinations of:
 - relying on private network
 - authenticating nodes using central authority
 - secure channels — no understanding without a key
 - combating isolated internal attacks (eg by voting techniques)
 - analysing any incoming mobile code
- beware: for some attacks no known cure

Heterogeneity support

- ? is a DS deployable through a variety of:
 - network types
 - hardware platforms + operating systems



Openness

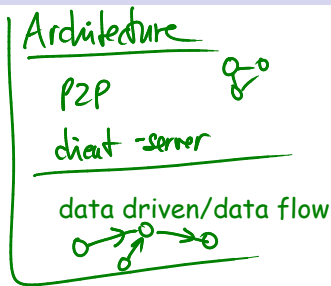
- how easy is it to extend and adapt a DS?
- a DS is *open* to the extreme if
 - all its protocols and APIs are public and well-documented
 - it provides well-documented access methods to all its resources
 - it is well-designed for extensibility, eg:
 - good decomposition (low coupling, cohesion, layers of abstraction)
 - no unnecessary complexity
 - it supports heterogeneity of networks, OSs and hardware
 - it is extensible in any common programming language

Transparency (1/2)

- transparency = certain aspect is not detectable (a tangible measure of abstraction)
- what to hide?
 - *access tr.*: local or remote?
 - *location tr.*: location in the net
 - *mobility tr.*: is it in a fixed location or moving?
 - *failure tr.*: failures of node hw/sw or communication hw/sw

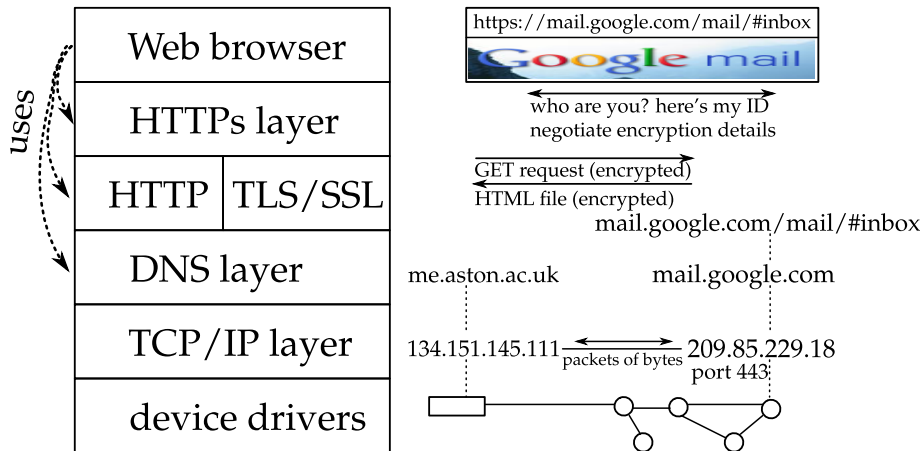
Transparency (2/2)

- what to hide?
 - *performance tr.*: overall load of DS
 - *scaling tr.*: size and power of DS
 - *replication tr.*: using multiple copies of a single resource
 - *concurrency tr.*: doing something concurrently



Network abstraction

Stacking abstraction layers



- network = nodes + links + communication facilities

Role of protocols

- *protocol* is a specification of
 - message formats
 - valid sequences of messages
- that allow multiple nodes cooperate to achieve a goal
- protocol usually defined in terms of lower-level protocol(s) layers of protocols
- abstract networks defined by a set of protocols
- eg Internet =
 - TCP/IP: reliable messaging + numerical addressing of nodes
 - DNS: hierarchical and human-friendly naming of nodes

Internet protocols for DS developers

- TCP provides logical channels between pairs of remote nodes
 - as if they were directly linked
 - can stream sequences of bytes of any length either way
 - hiding IP packets from users
- establishing a TCP channel:
 - one computer listens, the other initiates
 - multiple listeners → listening on *ports* 1–
- DNS lesser-known facts:
 - several DNS names can share one IP
 - one DNS name can refer to multiple computers
 - DNS-IP mapping sensitive to location

Web layer — resources and HTTP

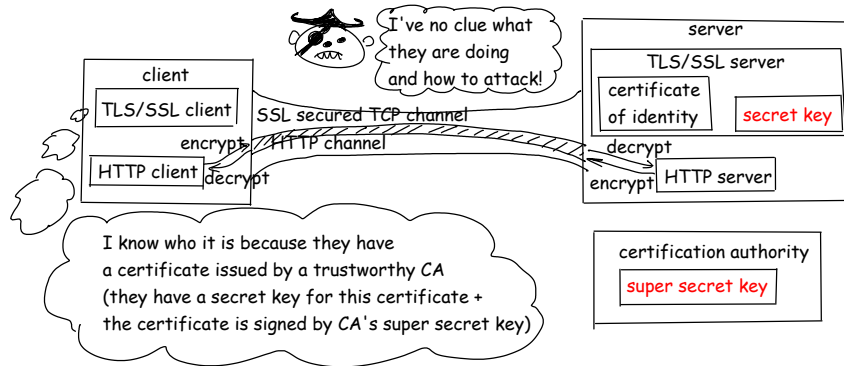
- resources = passive or active entities accessible over Internet
- can be viewed as logical nodes/sub-nodes for a DS
- URL — resource locator `http://a.com/d.pdf?page=1#fig1`

scheme, authority (DNS name + TCP port), path, query, fragment

- URI = URL or URN, URN has no DNS or TCP information, only a name, eg `urn:isbn:0-486-27557-4`
- HTTP: a request-response protocol for manipulating resources
- most common requests and their purpose:
 - GET: obtain a remote resource
 - POST: initiate operation on a remote resource
 - PUT: store a local resource at a remote location
 - DELETE: permanently remove a remote resource

Web layer — HTTPS

- HTTP over a secure channel:




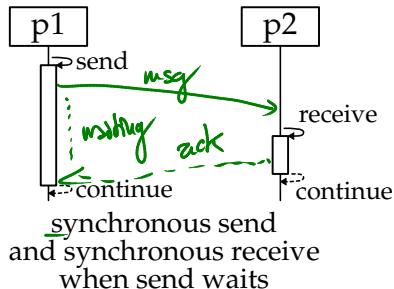


- protocol based on HTTP, TLS/SSL, TCP/IP, DNS

Network abstraction





Middleware — network abstraction

- characterised by
 - addressing logical nodes, not physical nodes
 - eg processes, resources, objects
 - results in location & access transparency
- convenient messaging, transporting typed structured data
 - types of data declared
 - greater openness and PL heterogeneity
 - synchronisation of communicating nodes

Synchronisation of communicating nodes




Synchronisation of communicating nodes




Multicast

- most efficient when implemented at TCP/IP level
- sometimes available at middleware level




naive multicast




efficient multicast

Multicast

- most efficient when implemented at TCP/IP level
- sometimes available at middleware level



naive multicast



efficient multicast

Remote administration

Why remote administration?

- install/update/configure/deploy a node's software
- read logs local to a node to diagnose a problem
- tweak some data files to fix a problem or simulate a problem
- monitor a node's load to tune a DS algorithm

Remote shell

- command line is simple but powerful — easy remote operation
- SSH: security achieved via secured TCP channels
- demo: install prg, view load, search files, rename files
- advantages/disadvantages?

Remote execution of GUI applications

- X Windows
 - thin-client: very efficient remote graphical clients
 - can have secure connection via SSH
 - not for MS Windows applications, only X-Windows applications
- Remote desktop
 - desktop completely taken over by remote computer
 - powerful but slow unless on very fast network
 - works on all OS, common in MS Windows administration

Learning Outcomes. You should now be able to

- define a DS
- name, describe and justify DS characteristics
- name and define important DS classifying aspects
- evaluate a given well-known DS against a set of classifying aspects
- list several levels of network abstraction behind a typical DS; for each level identify the key concepts that are relevant to a DS developer
- define and correctly use terms: TCP port, TCP channel, URL, secure channel, multicast
- explain difference between synchronous and asynchronous send
- give the main idea of an efficient multicast implementation
- start remote programs, copy files using SSH